

Homework 2 solution

```
using Random
using LinearAlgebra
using Convex
using ECOS
using Plots
using JLD
using CSV
using DataFrames
```

1. Since both $\mathbf{null}(A)$ and $\mathbf{null}(L)$ are subspaces, if $0 \neq u \in \mathbf{null}(A) \cap \mathbf{null}(L)$ then any scaling $\gamma u \in \mathbf{null}(A) \cap \mathbf{null}(L)$. So, if x^* solves the problem, then any $x^* + \gamma u$ also solves the problem, since $A(x^* + \gamma u) = Ax^*$ and $L(x^* + \gamma u) = Lx^*$.

To show the other direction, suppose that the solution to the optimization problem is not unique. That is, there exists $x \neq y$ where

$$(A^T A + \lambda L^T L)x = A^T b \text{ and } (A^T A + \lambda L^T L)y = A^T b.$$

Since we put no limitations on x and y (except that $x \neq y$) this implies that there exists some $u \neq 0$ where

$$(A^T A + \lambda L^T L)u = 0.$$

This implies

$$u^T A^T A u + \lambda u^T L^T L u = \|Au\|_2^2 + \lambda \|Lu\|_2^2 = 0.$$

Since both terms being added are nonnegative, both terms must be 0.

$$\|Au\|_2 = 0 \iff Au = 0, \quad \|Lu\|_2 = 0 \iff Lu = 0.$$


Therefore $u \in \mathbf{null}(A) \cap \mathbf{null}(L)$, and $u \neq 0$.

2. Multiobjective problems

a. **2-norm regularization** The solution is $x = (A^T A + \gamma I)^{-1} A^T b$. So,

$$\begin{aligned} \|x\|_2^2 &= \|(A^T A + \gamma I)^{-1} A^T b\|_2^2 \\ &= \|(Q D Q^T + \lambda Q Q^T)^{-1} A^T b\|_2^2 \\ &= \|(D - \lambda I)^{-1} Q^T A^T b\|_2^2 \\ &= \sum_{i=1}^n \frac{1}{(d_i + \gamma)^2} g_i^2. \end{aligned}$$

Sketch:



b. **Sparsity** The following code provides accuracy and sparsity measure for $\lambda = 1$.

```
#part 2a.
# load data
A = load("hw2_p2_smooth_A.jld")["data"]
b = load("hw2_p2_smooth_b.jld")["data"]
# load signal
x0 = load("hw2_p2_smooth_signal.jld")["data"]

(m,n) = size(A)

x_var = Variable(n);
```

```

loss = sumsquares(A*x_var-b)
reg = sum(abs(x_var[2:end]-x_var[1:end-1]))

λ = 1
problem = minimize(loss+λ*reg)
solve!(problem, ECOSolver(verbose=false))
xa = x_var.value;
free!(x_var);#
println("accuracy f1 = $(0.5*norm(A*xa-b,2)^2) and
sparsity f2 = $(norm(xa,1))")

```

The value of $f_1(x) = 0.619$ and $f_2(x) = 3.332$.

c. The following code plots the solution for $\lambda = 0.1$ and $\lambda = 10$.

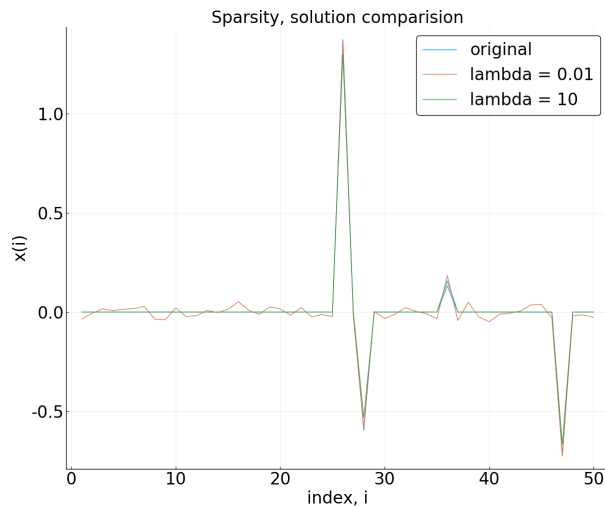
```

λ = .001
problem = minimize(loss+λ*reg)
solve!(problem, ECOSolver(verbose=false))
xb = x_var.value;
free!(x_var)

λ = 10
problem = minimize(loss+λ*reg)
solve!(problem, ECOSolver(verbose=false))
xc = x_var.value;
free!(x_var)

pyplot(size=(1200,1000), legend = true, legendfontsize=24,
xguidefontsize=24, yguidefontsize=24,
xtickfontsize=24, ytickfontsize=24, titlefontsize = 24)
plot_temp = plot(1:n, [x0,xb,xc], label = ["original"
"lambda = 0.01" "lambda = 10"], xlabel="index, i",
ylabel="x(i)", title = "Sparsity, solution comparision")
savefig(plot_temp, "figures/hw2_p2_c.png")

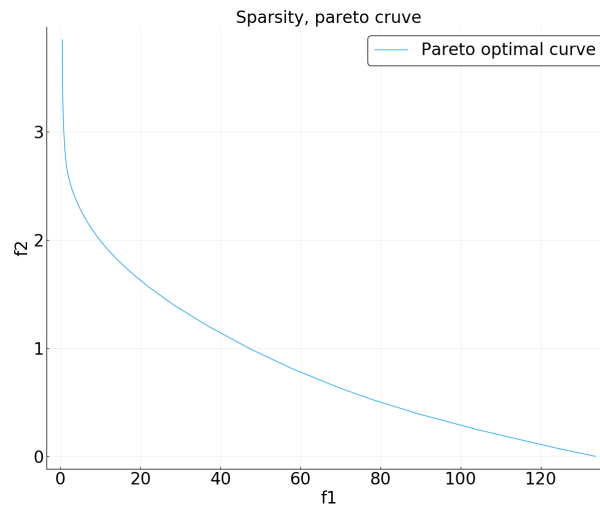
```



d. The following code plots the pareto optimal curve for the sparsity case.

```
# generate logspace between -3,3
Lambda = exp10.(range(-3, stop=3, length=100))
f1_values = []
f2_values = []
for i in 1:100
    λ = Lambda[i]
    problem = minimize(loss+λ*reg)
    solve!(problem, ECOSolver(verbose=0))
    xd = x_var.value;
    free!(x_var)
    push!(f1_values, 0.5*norm(A*xd-b,2)^2)
    push!(f2_values, norm(xd,1))
end

pyplot(size=(1200,1000), legend = true, legendfontsize=24,
xguidefontsize=24, yguidefontsize=24,
xtickfontsize=24, ytickfontsize=24, titlefontsize = 24)
plot_temp = plot(f1_values, f2_values, label = "Pareto
optimal curve", xlabel="f1", ylabel="f2", title = "Sparsity,
pareto cruve")
savefig(plot_temp, "figures/hw2_p2_d.png")
```



- e. The following code plots the solution for $\lambda = 0.1$ and $\lambda = 10$. It also provides the accuracy and total variation measure for $\lambda = 1$.

```
# load data
A = load("hw2_p2_smooth_A.jld")["data"]
b = load("hw2_p2_smooth_b.jld")["data"]

# load signal
x0 = load("hw2_p2_smooth_signal.jld")["data"]
x_var = Variable(n);

(m,n) = size(A)

loss = sumsquares(A*x_var-b)
reg = sum(abs(x_var[2:end]-x_var[1:end-1]))

λ = 1
problem = minimize(loss+λ*reg)
solve!(problem,ECOSSolver(verbose=false))
xa = x_var.value;
free!(x_var);#
println("accuracy f1 = $(0.5*norm(A*xa-b,2)^2) and
sparsity f2 = $(norm(xa,1))")

λ = .001
problem = minimize(loss+λ*reg)
solve!(problem,ECOSSolver(verbose=false))
xb = x_var.value;
free!(x_var)
```

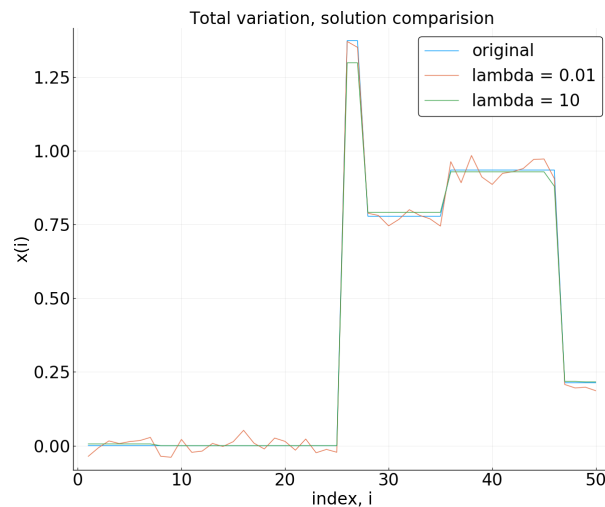
```

λ = 10
problem = minimize(loss+λ*reg)
solve!(problem, ECOSolver(verbose=false))
xc = x_var.value;
free!(x_var)

plot_temp = plot(1:n, [x0,xb,xc], label = ["original"
"lambda = 0.01" "lambda = 10"], xlabel="index, i",
ylabel="x(i)",
title = "Total variation, solution comparison")
savefig(plot_temp, "figures/hw2_p2_e.png")

```

For $\lambda = 1$, the value of $f_1(x) = 0.652$ and $f_2(x) = 20.310$.



f. The following code plots the pareto optimal curve for the total variation case.

```

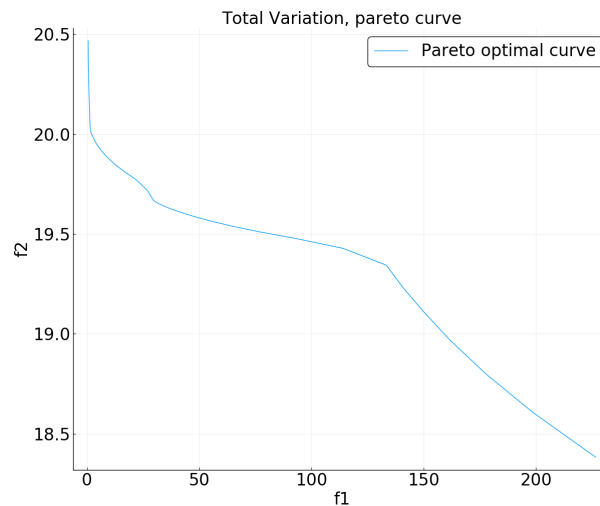
# generate logspace between -3,3
Lambda = exp10.(range(-3, stop=3, length=100))
f1_values = []
f2_values = []
for i in 1:100
    λ = Lambda[i]
    problem = minimize(loss+λ*reg)
    solve!(problem, ECOSolver(verbose=false))
    xd = x_var.value;
    free!(x_var)
    push!(f1_values, 0.5*norm(A*xd-b,2)^2)
    push!(f2_values, norm(xd,1))
end

```

```

end
plot_temp = plot(f1_values, f2_values, label = "Pareto
optimal curve", xlabel="f1", ylabel="f2", title = "Total
Variation, pareto curve")
savefig(plot_temp, "figures/hw2_p2_d.png")

```



3. Non-linear least squares

a. The gradient is

$$\nabla f(x) = 4 \sum_{i=1}^m (\|x - c_i\|_2^2 - d_i^2)(x - c_i).$$

$$\text{b. } r(x) = \begin{bmatrix} \|x - c_1\|_2^2 - d_1^2 \\ \vdots \\ \|x - c_m\|_2^2 - d_m^2 \end{bmatrix}, \quad J(x) = \begin{bmatrix} 2(x - c_1)^T \\ \vdots \\ 2(x - c_m)^T \end{bmatrix}$$

c. In the Gauss-Newton method, we attempt to solve at each iteration

$$\underset{x \in \mathbf{R}^n}{\text{minimize}} \quad \|r(\bar{x}) + J(\bar{x})(x - \bar{x})\|_2^2$$

where $\bar{x} = x^{(k)}$ the current iterate. Taking

$$A = J(\bar{x}) \in \mathbf{R}^{m \times n}, \quad b = J(\bar{x})\bar{x} - r(\bar{x})$$

the above problem becomes a least squares problem

$$\underset{x \in \mathbf{R}^n}{\text{minimize}} \quad \|Ax - b\|_2^2$$

with normal equations

$$A^T A x = A^T b.$$

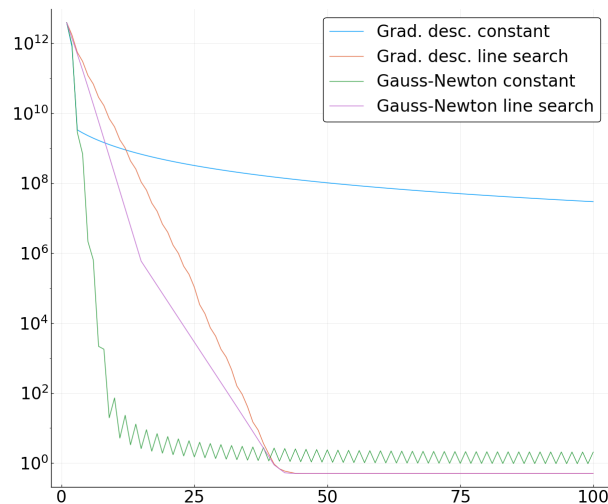
The solution to the normal equations is unique if and only if $A^T A$ is invertible—that is, if A has full column rank. Since we are assuming $n = 2$, we need only verify that $\text{rank}(A) = 2$.

Expanding

$$A = J(\bar{x}) = \begin{bmatrix} 2(\bar{x} - c_1)^T \\ \vdots \\ 2(\bar{x} - c_m)^T \end{bmatrix}$$

which has rank 2 under our assumptions.

d. The plot of $f(x^{(k)})$ for $k = 1, \dots, 100$, for all four solvers, is shown below:



Gradient descent: $\bar{\alpha} = 6.66 \times 10^{-8}$, whereas Gauss-Newton $\bar{\alpha} = 1$.
(Anything super small for GD and 1 for Gauss-Newton is acceptable.)

(Any answers for qualitative observations is acceptable.) Possible observations:

- i. Gauss Newton is much more tolerant of larger step sizes than gradient descent, suggesting better numerical stability.
- ii. Overall, using line search greatly improves method performance. However, there is a complexity tradeoff, since line search requires checking a condition many times.

Here's the code:

```
C = load("hw2_p3_C.jld")["data"]
d = load("hw2_p3_d.jld")["data"]
# load signal
x0 = load("hw2_p3_signal.jld")["data"]

s = 1
α = 0.5
β=0.5
epsilon = 10^(-4)

#Gradient descent with constant step size
x = []
obj_gd = []
push!(x, [1000;-500])

for iter in 1:100
    x_temp = x[iter]
    r = sum((C.-x_temp).^2,dims=1)' - d.^2;
    push!(obj_gd, norm(r,2)^2/2)
    g = 4*sum((ones(2,1)*r').*(x_temp.-C),dims = 2)
    push!(x, x_temp - g/15000000);
end

pyplot(size=(1200,1000), legend = true, legendfontsize=24,
xguidefontsize=24, yguidefontsize=24,
xtickfontsize=24, ytickfontsize=24, titlefontsize = 24)
plot_temp = plot(obj_gd,yaxis=:log, label = "Grad. desc.
constant")

#Gradient descent with line search
x = []
obj_gd_line = []
push!(x, [1000;-500])
```

```

for iter in 1:100
    x_temp = x[iter]
    r = sum((C.-x_temp).^2,dims=1)' - d.^2;
    push!(obj_gd_line, norm(r,2)^2/2)
    g = 4*sum((ones(2,1)*r').*(x_temp.-C),dims = 2)

    global t = s
    global obj_cand = obj_gd_line[iter]

    while (obj_gd_line[iter] .- obj_cand .- α*t*(g'*g))[1]
< 1e-10
        global t = t*β
        xcand = x_temp .- t*g
        rcand = sum((C.-xcand).^2,dims=1)' .- d.^2;
        global obj_cand = norm(rcand,2)/2
    end
    push!(x, x_temp - t*g);
end
plot!(obj_gd_line,yaxis=:log, label = "Grad. desc. line
search")

#Gauss newton with constant step size
x = []
obj_gd_gauss = []
push!(x, [1000;-500])

for iter in 1:100
    x_temp = x[iter]
    r = sum((C.-x_temp).^2,dims=1)' - d.^2
    push!(obj_gd_gauss, norm(r,2)^2/2)
    J = (x_temp.-C)'
    z = J \ r
    push!(x, x_temp-z)
end
plot!(obj_gd_gauss,yaxis=:log, label = "Gauss-Newton
constant")

#Gauss newton with line search
x = []
obj_gd_gauss_line = []
push!(x, [1000;-500])

for iter in 1:100
    x_temp = x[iter]

```

```

    r = sum((C.-x_temp).^2,dims=1)' - d.^2
    push!(obj_gd_gauss_line, norm(r,2)^2/2)
    J = (x_temp.-C)'
    z = J \ r
    global t = s
    global obj_cand = norm(r,2)^2/2
    g = 4*sum((ones(2,1)*r').*(x_temp.-C),dims = 2)

    while (obj_gd_gauss_line[iter] .- obj_cand .- α*t*
(g'*z))[1] < 1e-10
        global t = t*β
        xcand = x_temp .- t*z
        rcand = sum((C.-xcand).^2,dims=1)' .- d.^2;
        global obj_cand = norm(rcand,2)/2
    end

    push!(x, x_temp-t*z)
end
plot!(obj_gd_gauss_line,yaxis=:log,label= "Gauss-Newton
line search")
savefig(plot_temp, "figures/hw2_p3_d.png")

```